

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
In [3]: import warnings

warnings.filterwarnings('ignore')
```

```
In [4]: data = 'A:/Live.csv'

df = pd.read_csv(data)
```

## Check shape of the dataset

```
In [5]: df.shape
```

```
Out[5]: (7050, 16)
```

We can see that there are 7050 instances and 16 attributes in the dataset. In the dataset description, it is given that there are 7051 instances and 12 attributes in the dataset.

So, we can infer that the first instance is the row header and there are 4 extra attributes in the dataset. Next, we should take a look at the dataset to gain more insight about it.

## Preview the dataset

```
In [6]: df.head()
```

```
Out[6]:
```

	status_id	status_type	status_published	num_reactions	num_c
0	246675545449582_1649696485147474	video	4/22/2018 6:00	529	
1	246675545449582_1649426988507757	photo	4/21/2018 22:45	150	
2	246675545449582_1648730588577397	video	4/21/2018 6:17	227	
3	246675545449582_1648576705259452	photo	4/21/2018 2:29	111	
4	246675545449582_1645700502213739	photo	4/18/2018 3:22	213	

# View summary of dataset

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7050 entries, 0 to 7049
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status_id             7050 non-null   object
1   status_type           7050 non-null   object
2   status_published      7050 non-null   object
3   num_reactions         7050 non-null   int64
4   num_comments         7050 non-null   int64
5   num_shares            7050 non-null   int64
6   num_likes             7050 non-null   int64
7   num_loves             7050 non-null   int64
8   num_wows              7050 non-null   int64
9   num_hahas             7050 non-null   int64
10  num_sads               7050 non-null   int64
11  num_angrys            7050 non-null   int64
12  Column1                0 non-null      float64
13  Column2                0 non-null      float64
14  Column3                0 non-null      float64
15  Column4                0 non-null      float64
dtypes: float64(4), int64(9), object(3)
memory usage: 881.4+ KB
```

# Check for missing values in dataset

```
In [8]: df.isnull().sum()
```

```
Out[8]: status_id             0
status_type             0
status_published        0
num_reactions           0
num_comments            0
num_shares              0
num_likes               0
num_loves               0
num_wows                0
num_hahas               0
num_sads                0
num_angrys              0
Column1                 7050
Column2                 7050
Column3                 7050
Column4                 7050
dtype: int64
```

We can see that there are 4 redundant columns in the dataset. We should drop them before proceeding further.

## Drop redundant columns

```
In [9]: df.drop(['Column1', 'Column2', 'Column3', 'Column4'], axis=1, inplace=True)
```

## Again view summary of dataset

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7050 entries, 0 to 7049
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status_id             7050 non-null   object
1   status_type           7050 non-null   object
2   status_published      7050 non-null   object
3   num_reactions         7050 non-null   int64
4   num_comments         7050 non-null   int64
5   num_shares            7050 non-null   int64
6   num_likes             7050 non-null   int64
7   num_loves             7050 non-null   int64
8   num_wows              7050 non-null   int64
9   num_hahas            7050 non-null   int64
10  num_sads              7050 non-null   int64
11  num_angrys            7050 non-null   int64
dtypes: int64(9), object(3)
memory usage: 661.1+ KB
```

Now, we can see that redundant columns have been removed from the dataset.

We can see that, there are 3 character variables (data type = object) and remaining 9 numerical variables (data type = int64).

## View the statistical summary of numerical variables

```
In [11]: df.describe()
```

Out[11]:

	num_reactions	num_comments	num_shares	num_likes	num_loves	num_wows
<b>count</b>	7050.000000	7050.000000	7050.000000	7050.000000	7050.000000	7050.000000
<b>mean</b>	230.117163	224.356028	40.022553	215.043121	12.728652	1.289362
<b>std</b>	462.625309	889.636820	131.599965	449.472357	39.972930	8.719650
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	17.000000	0.000000	0.000000	17.000000	0.000000	0.000000
<b>50%</b>	59.500000	4.000000	0.000000	58.000000	0.000000	0.000000
<b>75%</b>	219.000000	23.000000	4.000000	184.750000	3.000000	0.000000
<b>max</b>	4710.000000	20990.000000	3424.000000	4710.000000	657.000000	278.000000

There are 3 categorical variables in the dataset. I will explore them one by one.

## Explore status\_id variable

In [12]: `df['status_id'].unique()`

Out[12]: `array(['246675545449582_1649696485147474',  
'246675545449582_1649426988507757',  
'246675545449582_1648730588577397', ...,  
'1050855161656896_1060126464063099',  
'1050855161656896_1058663487542730',  
'1050855161656896_1050858841656528'], dtype=object)`

In [13]: `len(df['status_id'].unique())`

Out[13]: 6997

We can see that there are 6997 unique labels in the status\_id variable. The total number of instances in the dataset is 7050. So, it is approximately a unique identifier for each of the instances. Thus this is not a variable that we can use. Hence, I will drop it.

## Explore status\_published variable

In [14]: `df['status_published'].unique()`

Out[14]: `array(['4/22/2018 6:00', '4/21/2018 22:45', '4/21/2018 6:17', ...,  
'9/21/2016 23:03', '9/20/2016 0:43', '9/10/2016 10:30'],  
dtype=object)`

In [15]: `len(df['status_published'].unique())`

Out[15]: 6913

Again, we can see that there are 6913 unique labels in the status\_published variable. The total number of instances in the dataset is 7050. So, it is also approximately a unique identifier for each of the instances. Thus this is not a variable that we can use. Hence, I will drop it also.

## Explore status\_type variable

```
In [16]: df['status_type'].unique()
```

```
Out[16]: array(['video', 'photo', 'link', 'status'], dtype=object)
```

```
In [17]: len(df['status_type'].unique())
```

```
Out[17]: 4
```

## Drop status\_id and status\_published variable from the dataset

```
In [18]: df.drop(['status_id', 'status_published'], axis=1, inplace=True)
```

## View the summary of dataset again

```
In [19]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7050 entries, 0 to 7049
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   status_type     7050 non-null   object
1   num_reactions   7050 non-null   int64
2   num_comments    7050 non-null   int64
3   num_shares      7050 non-null   int64
4   num_likes       7050 non-null   int64
5   num_loves       7050 non-null   int64
6   num_wows        7050 non-null   int64
7   num_hahas       7050 non-null   int64
8   num_sads        7050 non-null   int64
9   num_angrys      7050 non-null   int64
dtypes: int64(9), object(1)
memory usage: 550.9+ KB
```

## Preview the dataset again

```
In [20]: df.head()
```

```
Out[20]:
```

	status_type	num_reactions	num_comments	num_shares	num_likes	num_loves	num_w
0	video	529	512	262	432	92	
1	photo	150	0	0	150	0	
2	video	227	236	57	204	21	
3	photo	111	0	0	111	0	
4	photo	213	0	0	204	9	

```
In [21]: X = df
y = df['status_type']
```

```
In [22]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()

X['status_type'] = le.fit_transform(X['status_type'])
y = le.transform(y)
```

## View the summary of X

```
In [23]: X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7050 entries, 0 to 7049
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   status_type     7050 non-null   int32
1   num_reactions   7050 non-null   int64
2   num_comments    7050 non-null   int64
3   num_shares      7050 non-null   int64
4   num_likes       7050 non-null   int64
5   num_loves       7050 non-null   int64
6   num_wows        7050 non-null   int64
7   num_hahas       7050 non-null   int64
8   num_sads        7050 non-null   int64
9   num_angrys      7050 non-null   int64
dtypes: int32(1), int64(9)
memory usage: 523.4 KB
```

## Preview the dataset X

```
In [24]: X.head()
```

Out[24]:

	status_type	num_reactions	num_comments	num_shares	num_likes	num_loves	num_w
0	3	529	512	262	432	92	
1	1	150	0	0	150	0	
2	3	227	236	57	204	21	
3	1	111	0	0	111	0	
4	1	213	0	0	204	9	

In [25]: `cols = X.columns`

In [26]:

```

from sklearn.preprocessing import MinMaxScaler

ms = MinMaxScaler()

X = ms.fit_transform(X)

```

In [27]: `X = pd.DataFrame(X, columns=[cols])`

In [28]: `X.head()`

Out[28]:

	status_type	num_reactions	num_comments	num_shares	num_likes	num_loves	num_w
0	1.000000	0.112314	0.024393	0.076519	0.091720	0.140030	0.010000
1	0.333333	0.031847	0.000000	0.000000	0.031847	0.000000	0.000000
2	1.000000	0.048195	0.011243	0.016647	0.043312	0.031963	0.000000
3	0.333333	0.023567	0.000000	0.000000	0.023567	0.000000	0.000000
4	0.333333	0.045223	0.000000	0.000000	0.043312	0.013699	0.000000

In [29]:

```

from sklearn.cluster import KMeans

kmeans = KMeans(n_clusters=2, random_state=0)

kmeans.fit(X)

```

Out[29]:

▼ KMeans

KMeans(n\_clusters=2, random\_state=0)

In [30]: `kmeans.cluster_centers_`

```
Out[30]: array([[3.28506857e-01, 3.90710874e-02, 7.54854864e-04, 7.53667113e-04,
 3.85438884e-02, 2.17448568e-03, 2.43721364e-03, 1.20039760e-03,
 2.75348016e-03, 1.45313276e-03],
 [9.54921576e-01, 6.46330441e-02, 2.67028654e-02, 2.93171709e-02,
 5.71231462e-02, 4.71007076e-02, 8.18581889e-03, 9.65207685e-03,
 8.04219428e-03, 7.19501847e-03]])
```

The KMeans algorithm clusters data by trying to separate samples in  $n$  groups of equal variances, minimizing a criterion known as inertia, or within-cluster sum-of-squares. Inertia, or the within-cluster sum of squares criterion, can be recognized as a measure of how internally coherent clusters are. The k-means algorithm divides a set of  $N$  samples  $X$  into  $K$  disjoint clusters  $C$ , each described by the mean  $\mu_j$  of the samples in the cluster. The means are commonly called the cluster centroids. The K-means algorithm aims to choose centroids that minimize the inertia, or within-cluster sum of squared criterion.

## Inertia

Inertia is not a normalized metric.

The lower values of inertia are better and zero is optimal.

But in very high-dimensional spaces, euclidean distances tend to become inflated (this is an instance of curse of dimensionality).

Running a dimensionality reduction algorithm such as PCA prior to k-means clustering can alleviate this problem and speed up the computations.

We can calculate model inertia as follows:-

```
In [31]: kmeans.inertia_
```

```
Out[31]: 237.7572640441956
```

The lesser the model inertia, the better the model fit.

We can see that the model has very high inertia. So, this is not a good model fit to the data.

```
In [32]: labels = kmeans.labels_

# check how many of the samples were correctly labeled
correct_labels = sum(y == labels)

print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
```

```
Result: 63 out of 7050 samples were correctly labeled.
```


```
In [33]: print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))
```

```
Accuracy score: 0.01
```



We have achieved a weak classification accuracy of 1% by our unsupervised model.

```
In [34]: from sklearn.cluster import KMeans
cs = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters = i, init = 'k-means++', max_iter = 300, n_init = 10)
    kmeans.fit(X)
    cs.append(kmeans.inertia_)
plt.plot(range(1, 11), cs)
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('CS')
plt.show()
```



By the above plot, we can see that there is a kink at  $k=2$ .

Hence  $k=2$  can be considered a good number of the cluster to cluster this data.

But, we have seen that I have achieved a weak classification accuracy of 1% with  $k=2$ .

I will write the required code with  $k=2$  again for convinience.

```
In [35]: from sklearn.cluster import KMeans

kmeans = KMeans(n_clusters=2, random_state=0)

kmeans.fit(X)
```

```

labels = kmeans.labels_

# check how many of the samples were correctly labeled

correct_labels = sum(y == labels)

print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))

print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))

```

Result: 63 out of 7050 samples were correctly labeled.  
Accuracy score: 0.01

So, our weak unsupervised classification model achieved a very weak classification accuracy of 1%.

I will check the model accuracy with different number of clusters.

## K-Means model with 3 clusters

```

In [36]: kmeans = KMeans(n_clusters=3, random_state=0)

kmeans.fit(X)

# check how many of the samples were correctly labeled
labels = kmeans.labels_

correct_labels = sum(y == labels)
print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))

```

Result: 138 out of 7050 samples were correctly labeled.  
Accuracy score: 0.02

## K-Means model with 4 clusters

```

In [37]: kmeans = KMeans(n_clusters=4, random_state=0)

kmeans.fit(X)

# check how many of the samples were correctly labeled
labels = kmeans.labels_

correct_labels = sum(y == labels)
print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))

```

Result: 4340 out of 7050 samples were correctly labeled.  
Accuracy score: 0.62

We have achieved a relatively high accuracy of 62% with k=4.

I have implemented the most popular unsupervised clustering technique called K-Means Clustering.

I have applied the elbow method and find that  $k=2$  ( $k$  is number of clusters) can be considered a good number of cluster to cluster this data.

I have find that the model has very high inertia of 237.7572. So, this is not a good model fit to the data.

I have achieved a weak classification accuracy of 1% with  $k=2$  by our unsupervised model.

So, I have changed the value of  $k$  and find relatively higher classification accuracy of 62% with  $k=4$ .

Hence, we can conclude that  $k=4$  being the optimal number of clusters.

In [ ]: